

## README.md

# Käyttöjärjestelmät ja systeemiohjelmointi

---

Sourcecode available at: <https://github.com/JaakkoLipp/SYS-Projects>

screenshots of programs functionality in screenshots/ folder

Ideally the documentation would be in seperate files

## Project 1: Warmup to C and Unix programming

---

### Overview

---

`reverse.c` reads lines from an input file, stores them in a dynamic array, and then writes them to an output file in reverse order. If no output file is provided, it prints the lines to the standard output.

### Usage

---

Run the program in a linux terminal with the following syntax:

```
reverse <input> <output>
```

Where:

- `<input>` is the name of the input file.
- `<output>` is the name of the output file (optional).

### Error Handling

---

The program includes the following error handling:

- If the input file and output file are the same, it prints an error message "Input and output file must differ" and exits with return code 1.

- If the user specifies an input file or output file that cannot be opened, it prints an error message "error: cannot open file 'filename'" and exits with return code 1.
- If `malloc()` fails to allocate memory, it prints the error message "malloc failed" and exits with return code 1.
- If the user runs the program with too many arguments, it prints "usage: reverse " and exits with return code 1.

All error messages are printed to `stderr` rather than `stdout` .

## Code Structure

---

The program begins by checking the command-line arguments and opening the input file. It then reads lines from the input file into a dynamic array, resizing the array as needed.

If an output file is specified, it opens the output file for writing. It then writes the lines from the array to the output file in reverse order, or to `stdout` if no output file is specified.

Finally, it frees the memory allocated for the array and closes the output file if it is open.

## Memory Management

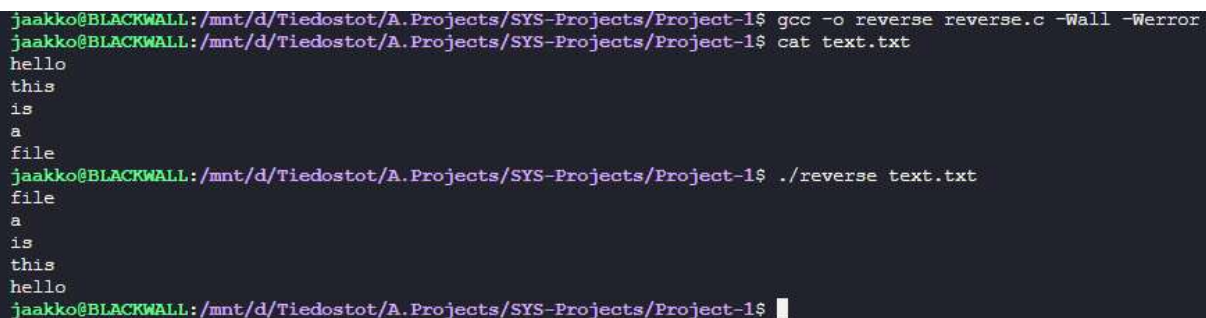
---

The program uses dynamic memory allocation to store the lines from the input file. It starts with an array of a certain size and doubles the size of the array whenever it becomes full. It also uses dynamic memory allocation to store each line of text, resizing the buffer as needed to accommodate long lines.

When the program has finished writing the lines to the output file, it frees the memory allocated for each line and for the array itself. This ensures that the program does not leak memory, even if the input file is very large.

## Screenshot

---



```
jaakko@BLACKWALL: /mnt/d/Tiedostot/A.Projects/SYS-Projects/Project-1$ gcc -o reverse reverse.c -Wall -Werror
jaakko@BLACKWALL: /mnt/d/Tiedostot/A.Projects/SYS-Projects/Project-1$ cat text.txt
hello
this
is
a
file
jaakko@BLACKWALL: /mnt/d/Tiedostot/A.Projects/SYS-Projects/Project-1$ ./reverse text.txt
file
a
is
this
hello
jaakko@BLACKWALL: /mnt/d/Tiedostot/A.Projects/SYS-Projects/Project-1$
```

# Project 2: Unix Utilities Page

---

## my-cat.c

---

### Overview

---

The `my-cat` program is a simple implementation similar to the Unix `cat` utility. It concatenates and displays the content of one or more files to the standard output. If no files are specified, the program exits without doing anything, effectively performing a no-operation (no-op).

### Usage

---

```
my-cat [file ...]
```

- `[file ...]` : Optional. One or more files to concatenate and display. If no files are provided, `my-cat` performs a no-op.

### Compilation

---

To compile the `my-cat` program, use the following `gcc` command:

```
gcc -o my-cat my-cat.c
```

This command compiles the source code into an executable named `my-cat`.

### Functionality

---

The main functionalities of the `my-cat` program include:

- **No Files Specified:** If invoked without any file arguments, the program exits immediately with a return code of `0`, performing no actions.
- **Reading and Displaying File Contents:** For each file specified in the command line arguments, the program attempts to open the file and read its content. The content is then printed to the standard output ( `stdout` ).
- **Error Handling:** If a file cannot be opened, an error message `my-cat: cannot open file` is printed to the standard output, and the program exits with a return code of `1`.

## Examples

---

To display the content of a single file:

```
./my-cat file1.txt
```

To concatenate and display the contents of multiple files:

```
./my-cat file1.txt file2.txt file3.txt
```

## Error Messages

---

- If the program fails to open a file, it prints `my-cat: cannot open file` to the standard output and exits with a return code of `1`.

## Exit Codes

---

- `0` : Program executed successfully, all specified files were opened, and their contents displayed.
- `1` : Program encountered an error, such as failure to open a specified file.

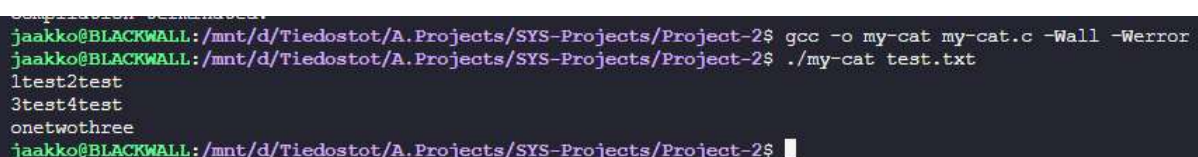
## Implementation Details

---

- The program iterates over each argument provided to it, treating each as a filename to be opened and read.
- For each file, it reads and prints the content using a buffer of fixed size ( `1024` bytes). This size is a compromise between memory usage and performance but might not be optimal for all situations.
- It uses the standard C library functions `fopen`, `fgets`, and `fclose` for file operations.
- In case a file cannot be opened, it prints an error message and exits with a non-zero status code, indicating failure.

## Screenshot

---



```
jaakko@BLACKWALL: /mnt/d/Tiedostot/A.Projects/SYS-Projects/Project-2$ gcc -o my-cat my-cat.c -Wall -Werror
jaakko@BLACKWALL: /mnt/d/Tiedostot/A.Projects/SYS-Projects/Project-2$ ./my-cat test.txt
1test2test
3test4test
onetwothree
jaakko@BLACKWALL: /mnt/d/Tiedostot/A.Projects/SYS-Projects/Project-2$
```

# my-grep.c

---

## Overview

---

The `my-grep` program is a simplified version of the Unix `grep` utility. It searches for lines containing a specified search term in one or more files and prints those lines to the standard output. If no file is provided, `my-grep` reads from standard input (stdin), allowing it to be used in a pipeline of commands.

## Usage

---

```
my-grep searchterm [file ...]
```

- `searchterm` : The string to search for within each file or stdin.
- `[file ...]` : Optional. One or more files to search. If no files are provided, `my-grep` searches stdin.

## Compilation

---

To compile the `my-grep` program, use the following command:

```
gcc -o my-grep my-grep.c
```

## Functionality

---

The program consists of two main parts: the `grep` function and the `main` function.

### grep Function

```
void grep(FILE *f, char *searchterm);
```

- **Purpose:** Searches a single file (or stdin) for lines containing the specified search term and prints those lines.
- **Parameters:**
  - `FILE *f` : A file pointer to the file being searched. Can also be stdin.
  - `char *searchterm` : The string to search for in the file's content.
- **Implementation Details:**
  - Reads the file line by line.

- Uses `strstr` to search for the `searchterm` in each line.
- If `searchterm` is found, the line is printed to `stdout`.

## Main Function

The `main` function processes command-line arguments, handles file opening and closing, and calls the `grep` function accordingly.

- **Command-Line Arguments:**
  - The first argument after the program name is considered the `searchterm`.
  - Any subsequent arguments are treated as filenames to be searched.
- **Behavior:**
  - If no arguments are provided, it prints usage information and exits.
  - If only `searchterm` is provided (i.e., `argc == 2`), it reads from `stdin`.
  - If one or more filenames are provided, it opens and searches each file for the `searchterm`.
- **Error Handling:**
  - If a file cannot be opened, it prints an error message and exits.

## Examples

---

Search for the term "error" in a file named "log.txt":

```
./my-grep error log.txt
```

Search for the term "foo" across multiple files:

```
./my-grep foo file1.txt file2.txt
```

Use in a pipeline to search for "hello" in the output of another command:

```
echo -e "hello\nworld" | ./my-grep hello
```

## Error Messages

---

- If the program is invoked without a search term, it will display: `my-grep: searchterm [file ...]` and exit.
- If a specified file cannot be opened, it will display: `my-grep: cannot open file and exit.`

## Exit Codes

---

- 0 : Program executed successfully.
- 1 : Program encountered an error, such as missing arguments or unable to open a file.

## Screenshot

---

```
jaakko@BLACKWALL: /mnt/d/Tiedostot/A.Projects/SYS-Projects/Project-2$ ./my-cat test.txt
1test2test
3test4test
onetwothree
jaakko@BLACKWALL: /mnt/d/Tiedostot/A.Projects/SYS-Projects/Project-2$ gcc -o my-grep my-grep.c -Wall -Werror
jaakko@BLACKWALL: /mnt/d/Tiedostot/A.Projects/SYS-Projects/Project-2$ ./my-grep 4 test.txt
3test4test
jaakko@BLACKWALL: /mnt/d/Tiedostot/A.Projects/SYS-Projects/Project-2$
```

# my-zip Documentation

---

## Overview

---

my-zip is a file compression tool that utilizes Run-Length Encoding (RLE) for compressing text files. It reads one or more text files as input and outputs a compressed version to the standard output (stdout), which is typically redirected to a file.

## Functionality

---

- **Input Handling:** my-zip takes one or more file names as command-line arguments. It reads each file in sequence.
- **Compression Algorithm:** The tool reads characters from the input stream and counts consecutive occurrences of the same character. For each sequence of identical characters, it outputs a 4-byte integer representing the count, followed by the character itself. This data is written in binary format to stdout.
- **Output:** The compressed data is outputted as binary data. Each entry consists of a 4-byte integer (count) followed by a single character, representing the run-length encoding of the input files.

## Usage

---

```
my-zip <file1> [file2 ...] > output.z
```

- Replace <file1> [file2 ...] with the names of the files you wish to compress.
- The output is binary and should be redirected to a file (e.g., output.z ).

## Error Handling

---

- If no files are provided, `my-zip` prints a usage message and exits with a status code of 1.
- If a file cannot be opened, an error message is printed, and `my-zip` exits with a status code of 1.

# my-unzip Documentation

---

## Overview

---

`my-unzip` is a file decompression tool designed to reverse the compression applied by `my-zip`. It reads a compressed file in the specific binary format produced by `my-zip` and outputs the original uncompressed text to standard output (stdout).

## Functionality

---

- **Input Handling:** `my-unzip` takes one or more compressed files as command-line arguments. It reads each file in sequence.
- **Decompression Algorithm:** The tool reads 5-byte entries from the input stream. Each entry consists of a 4-byte integer (count) and a single character. `my-unzip` then prints the character `count` times to stdout, effectively decompressing the data.
- **Output:** The decompressed text is outputted to stdout and can be redirected to a file or another command.

## Usage

---

```
my-unzip input.z > output.txt
```

- Replace `input.z` with the name of the compressed file you wish to decompress.
- The output is text and can be redirected to a file (e.g., `output.txt`).

## Error Handling

---

- If no files are provided, `my-unzip` prints a usage message and exits with a status code of 1.
- If a file cannot be opened, an error message is printed, and `my-unzip` exits with a status code of 1.



## Screenshot

```
jaakko@BLACKWALL: /mnt/d/Tiedostot/A.Projects/SYS-Projects/Project-2$ gcc -o my-unzip my-unzip.c -Wall -Werror
jaakko@BLACKWALL: /mnt/d/Tiedostot/A.Projects/SYS-Projects/Project-2$ gcc -o my-zip my-zip.c -Wall -Werror
jaakko@BLACKWALL: /mnt/d/Tiedostot/A.Projects/SYS-Projects/Project-2$ cat compressiontest
aaaaaaaaabbbb
jaakko@BLACKWALL: /mnt/d/Tiedostot/A.Projects/SYS-Projects/Project-2$ ./my-zip compressiontest > compressed
jaakko@BLACKWALL: /mnt/d/Tiedostot/A.Projects/SYS-Projects/Project-2$ ./my-unzip compressed
aaaaaaaaabbbb
jaakko@BLACKWALL: /mnt/d/Tiedostot/A.Projects/SYS-Projects/Project-2$
```

# Project 4: Modified XV6 Kernel with New Syscall `getreadcount`

## Overview

This project introduces a modified XV6 kernel that implements a new syscall, `getreadcount`. This syscall returns the number of times the `read()` system call has been invoked by user processes since the kernel was booted.

## Location

The modified kernel can be explored and tested within the `kernel` folder inside Project 4.

## Purpose

The purpose of adding the `getreadcount` system call is to monitor the usage of the `read()` system call, providing insights into read operations performed by user processes.

## Implementation Summary

### Adding `getreadcount` System Call to XV6

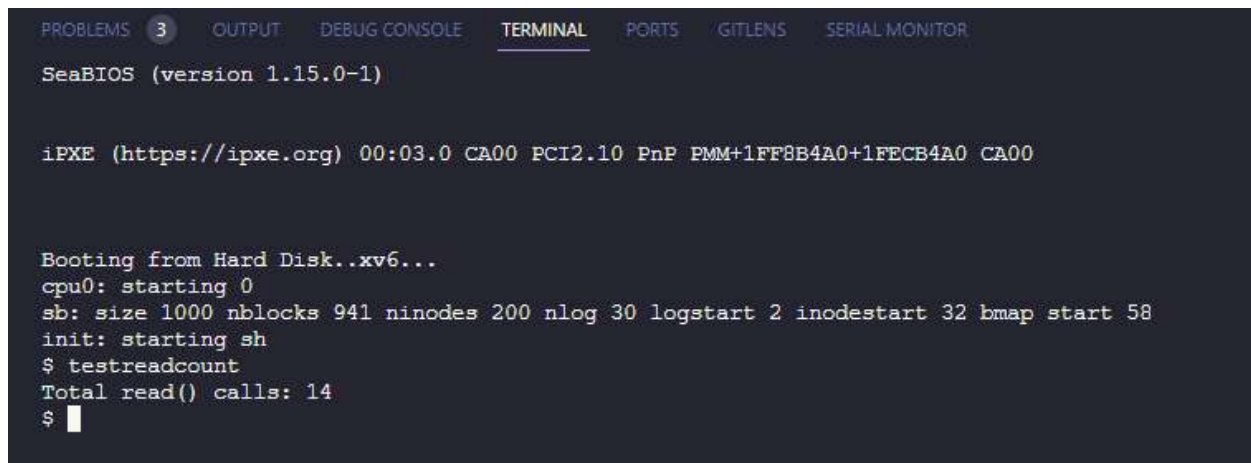
- **Global Counter:** Introduced a global variable `uint readcount` to keep track of `read()` invocations.
- **Modifying `read` Function:** The `readcount` is incremented within the `read()` function to reflect each invocation.
- **System Call Implementation:** Implemented the `sys_getreadcount` function to return the current value of `readcount`.

- **System Call Registration:** Assigned a unique identifier to the new system call and updated the syscall table in `syscall.c` to include it.
- **User Space Accessibility:** Made `getreadcount()` accessible to user programs by declaring it in `user.h`.

## Testing

- Developed a test program `testreadcount` to execute several `read()` operations followed by a `getreadcount()` call. The program then prints the result to verify the accurate incrementation of the counter.
- The project can be compiled using `make clean` followed by `make` or `make qemu-nox` for execution (Note: `qemu` is required).
- Confirmed to function correctly when tested in a `qemu /Ubuntu` host environment.

## Screenshot

A screenshot of a terminal window with a dark background and light-colored text. The terminal shows the output of a program running in a QEMU environment. At the top, there are tabs for 'PROBLEMS', 'OUTPUT', 'DEBUG CONSOLE', 'TERMINAL' (which is selected), 'PORTS', 'GIT LENS', and 'SERIAL MONITOR'. The output text includes: 'SeaBIOS (version 1.15.0-1)', 'iPXE (https://ipxe.org) 00:03.0 CA00 PCI2.10 PnP PMM+1FF8B4A0+1FECB4A0 CA00', 'Booting from Hard Disk..xv6...', 'cpu0: starting 0', 'sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58', 'init: starting sh', '\$ testreadcount', 'Total read() calls: 14', and '\$ ' followed by a cursor.

This screenshot showcases the output of the `testreadcount` program, demonstrating the functionality of the new `getreadcount` syscall.

### disclaimers and sources

All programs have been tested to compile and work with `-Wall -Werror` in `gcc`. Generative AI was used to generate and improve parts of the documentation, and to refine and improve code readability and structure. AI was used as a teaching tool and some information may be based on AI generated content.